



**NOAA**  
**FISHERIES**

# Using Deep Learning for DNA classification tasks


**Chris Powers**

SEDNA Office Hours

20 November 2025

# Machine learning and deep learning as statistical frameworks



- **Artificial intelligence (AI)** is the broadest framework, defined as creating so-called “intelligent machines”
- **Machine learning (ML)**, a subset of AI, allows intelligent machines to identify patterns and trends without explicit direction
- **Deep learning (DL)**, a subset of ML, allows for the handling of complex tasks and unstructured data



**NOAA**  
FISHERIES


# ML may seem inaccessible, but most of us use it in science without realizing

- ML is broadly considered to be highly complex, but this is not necessarily the case
- Simple examples of ML are part of every biologist's statistical toolkit:  
Simple linear regression  
Generalized additive models  
Classification trees  
Maximum likelihood analysis  
Clustering
- Today, deep-learning is the focus  
Convolutional neural networks



**NOAA**  
FISHERIES

# Deep learning models often rely on a neural network structure




[https://tikz.net/neural\\_networks/](https://tikz.net/neural_networks/)

- Unstructured input data in some tensor format
- Tunable parameters in “hidden layers”
- Output layer relatable to question asked of model (e.g. number of classification labels)
- Number of hidden layers is proportional to capacity to learn
- Iteratively **tune** parameters over many **epochs**, while monitoring accuracy and confidence



**NOAA**  
FISHERIES

# Deep-learning models iterate to find local minima




- Apply gradient descent (or something like it) to minimize errors  
Note, actual framework is an nth dimensional space
- Every step represents iteratively updated weights on neurons in a network
- Once the model stops improving, or reaches a minima, the model is done training

<https://blog.paperspace.com/part-3-generic-python-implementation-of-gradient-descent-for-nn-optimization/>

# Deep-learning models require best practices in data analysis, including training/validation splits

- Deep learning models have an enormous capacity to learn!
- Correlates with a large capacity to **memorize**
- The perfect end in a gradient descent would be memorization of the dataset
- Ultimate goal: learning trends in underlying data to **generalize**
- How do we achieve this?  
Assess performance on unseen data, not training data



**NOAA**  
**FISHERIES**




**NOAA**  
**FISHERIES**

# Convolutional Neural Networks

# Convolutional neural networks (CNNs) extract abstract features that are useful for classification (Computer vision)

A convolutional neural network is an older type of deep learning model (1990s), designed for processing grid-like data such as images, that learns features by applying a set of filters or kernels to the input data





Example of image classification convolution  
From:  
<https://i.sstatic.net/bN2iA.png>




**NOAA**  
FISHERIES

# Convolution functions apply a sliding window approach to transform the data

- Example: Slide a kernel (or filter) over an array of data and perform a dot product multiplication with the kernel  
Extracts unique features of the kernel
- Isolates a local signal across an image in the hidden layer part of the network
- This can train the neural network on emergent trends




2-Dimensional convolution. From:  
[https://ujjwalkarn.me/wp-content/uploads/2016/07/convolution\\_schematic.gif?w=268&h=196](https://ujjwalkarn.me/wp-content/uploads/2016/07/convolution_schematic.gif?w=268&h=196)



[https://tikz.net/neural\\_networks/](https://tikz.net/neural_networks/)

# Why use an image classification technique for DNA?


- CNNs first used on DNA in 2015
- DNA sequences contain a wide variety of conserved motifs and domains that encode functional and taxonomic information
- The convolutional feature extraction process isolates motifs and domains



CNN architecture of a DNA model,  
DL-TODA  
From: Cres et. al. 2023,  
*Biomolecules*



**NOAA**  
FISHERIES



## Convolutional Neural network approach out-performs gold-standard tools for taxonomic assignment at species level

- The convolutions enable the model to pick up on obscure, hidden signals

CNN architecture of a DNA model,  
DL-TODA  
From: Cres et. al. 2023,  
*Biomolecules*




**NOAA**  
**FISHERIES**

# To prep DNA for classification, it needs to be embedded into a convolvable tensor

- DNA sequences are not inherently readable by convolutional neural networks
  - Convert the data to a machine-readable tensor
- Common techniques include:
  - One-hot encoding
  - K-mer tokenization

A C G A G G T A C G A



1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x2</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature

2-Dimensional convolution

From:

[https://ujjwalkarn.me/wp-content/uploads/2016/07/convolution\\_schematic.gif?w=268&h=196](https://ujjwalkarn.me/wp-content/uploads/2016/07/convolution_schematic.gif?w=268&h=196)





**NOAA**  
FISHERIES

# To prep DNA for classification, it needs to be embedded into a convolvable vector

- DNA sequences are not inherently readable by convolutional neural networks
  - Convert the data to a machine-readable vector
- Common techniques include:
  - One-hot encoding
  - K-mer tokenization

A	T	G	C	-	A
1	0	0	0	0	1
0	1	0	0	0	0
0	0	1	0	0	0
0	0	0	1	0	0





[https://tikz.net/neural\\_networks/](https://tikz.net/neural_networks/)

- The input vector is defined as the green input layer
- The hidden layers are defined by the size of the model  
The convolutions happen in hidden layers
- The output layers are defined by the **task**




**NOAA**  
**FISHERIES**

# Defining a meaningful ML task for the CNN

- The output layer will be a **vector of probabilities** associated with a label

**Taxonomic Prediction**




**Taxa 1**


**Taxa 2**

**Taxa 3**

**Gene Calling**




**Protein Function Prediction**



The model trains iteratively to maximize accuracy (minimize loss) over  $n$  epochs

# Overall: CNNs offer a lot of flexibility to deeply analyze trends in DNA

- Older technique, but reliable
- Visualize hidden trends through convolutions
- Classification output can be highly flexible
- Complexity of the tasks correlates with the size of the network correlates with capacity to learn
- **Questions on CNNs?**



[https://tikz.net/neural\\_networks/](https://tikz.net/neural_networks/)




**NOAA**  
**FISHERIES**

# Case study: Using deep-learning to improve DNA classification

# Deep learning and fish eDNA metabarcoding

- Environmental DNA (eDNA) can be used to target a specific taxonomic group for sequencing based on shed DNA
- Fish eDNA is commonly sequenced using two primers in the 12S region
- Many species share the same sequence, limiting taxonomic resolution
- Current sequence classification approaches operate near theoretical maximum accuracy
- Other information may be used to identify these sequences




# ePlacer: Using Deep-Learning to augment taxonomic assignment.

- Use CNN to set up a DNA classification network
- Add a geographic location tensor From OBIS
- Combine the geographic and DNA vectors
- Resolve fish with exact sequence overlap, but distinct geography
- Here, models were trained and tested for 2 fish eDNA marker genes

A	T	G	C	-	A
1	0	0	0	0	1
0	1	0	0	0	0
0	0	1	0	0	0
0	0	0	1	0	0

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0.1	0.5	0.1	0	0	0	0
0	0	0	0.5	1	0.5	0	0	0	0
0	0	0	0.1	0.5	0.1	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0





[https://tikz.net/neural\\_networks/](https://tikz.net/neural_networks/)




**NOAA**  
FISHERIES

# ePlacer: Using Deep-Learning to augment taxonomic assignment.


A	T	G	C	-	A
1	0	0	0	0	1
0	1	0	0	0	0
0	0	1	0	0	0
0	0	0	1	0	0

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0.1	0.5	0.1	0	0	0	0	0
0	0	0	0.5	1	0.5	0	0	0	0	0
0	0	0	0.1	0.5	0.1	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0




NOAA  
FISHERIES


# Both marker gene regions exhibited robust learning patterns



**NOAA**  
FISHERIES




**NOAA**  
**FISHERIES**




**NOAA**  
**FISHERIES**

MiFish in situ



Riaz in situ




- Comparison of the geographically aware CNN to gold-standard classification
- Dataset consists of real-world data from Stoeckle *et. al.* 2021 and Ledger *et. al.* In Prep.
- ePlacer exhibited over-performance of gold standard, sequence-only tools across all performance metrics


Tool    ePlacer Curated    ePlacer    Qiime2    Kraken2



NOAA  
FISHERIES




- Comparison of the geographically aware CNN to gold-standard classification
- Dataset consists of real-world data from Stoeckle *et. al.* 2021 and Ledger *et. al.* In Prep.
- ePlacer exhibited over-performance of gold standard, sequence-only tools across all performance metrics



**NOAA**  
**FISHERIES**

# A practical application on MiFish: Examination of the *Clupea*



```
CGCGTAAAGAGTGGTTATGGAAAACAAGCACTAA  
AGCCAAAGAGCCCTCAGGCCGTATACGCACCCG  
GGGCCTCGAACCACTATCACGAAAGTAGCTTTACC  
CTCGCCCACCAGAACCCACGAGAGCTGGGACACA
```

The *Clupea* are difficult to discriminate



*Clupea pallasii*  
Manual curation

*Clupea*  
Qiime: 0.999 Confidence




**NOAA**  
FISHERIES


# A practical application on MiFish: Examination of the *Clupea*

```
CGCGTAAAGAGTGGTTATGGAA  
ACAAAGCACTAAAGCCAAAGAGC  
CCTCAGGCCGTTATACGCACCCG  
GGGCCTCGAACCACTATCACGAA  
AGTAGCTTACCCTGCCACCAG  
AACCACGAGAGCTGGGACACA
```

ePlacer is biogeographically aware,  
**increasing confidence of**  
assignment for *Clupea pallasii*



*Clupea pallasii*  
ePlacer: 0.999 Confidence



**NOAA**  
**FISHERIES**

# A practical application on MiFish: Examination of the *Clupea*



```
CGGCGTAAAGAGTGGTTATGGAA  
AACAAAGCACTAAAGCCAAAGAGC  
CCTCAGGCCGTATACGCACCCG  
GGGCCTCGAACCACTATCACGAA  
AGTAGCTTACCCTCGCCCACCAAG  
AACCCACGAGAGCTGGGACACA
```



The ability to use the same tool and databases for taxonomic assignment increases **Interoperability** and **FAIRness** of results



*Clupea pallasi*  
ePlacer: 0.999 Confidence



*Clupea harengus*  
ePlacer: 0.784 Confidence



**NOAA**  
**FISHERIES**

# Augmenting DNA classification using Deep-Learning approaches


- Deep-learning approaches are powerful when applied to DNA-related problems
- Example: ePlacer substantially increases the accuracy and interoperability of taxonomic assignment
- **Future Directions**
- Incorporate additional data modalities to estimate impacts of ecological frequency in addition to biogeography
- Depth, seasonality, water temperature, distance to coast
- **Questions?**



**NOAA**  
**FISHERIES**

# Monitoring performance of deep-learning: Some key terms

- **Epoch:** An iteration of training. In an epoch, the model has seen the whole training dataset.
- **Training Dataset:** Data used to tune hidden layer parameters. Typically ~ 70%
- **Validation Dataset:** Data used to check performance of model. Typically ~ 30%
- **Accuracy:** Total proportion of correct classifications (correctness)
- **Loss:** A measure of the distance between predicted and actual classification (confidence)



**NOAA**  
FISHERIES



**NOAA**  
**FISHERIES**

# Transformer-Based Models

# For transformer models, all you need is attention

- Cutting edge technology
- No reliance on transformation/convolution
- Transformer-based models use position-based encodings and attention to learn patterns across the whole sequence





Figure 1: The Transformer - model architecture.



From: Vaswani et. al. 2017




From:  
<https://www.exxactcorp.com/blog/Deep-Learning/a-deep-dive-into-the-transformer-architecture-the-development-of-transformer-models>

# Why are transformer models useful for DNA?

- Transformer models have been best applied for natural language processing (NLP) tasks
- Current NLP applications are easily adaptable to DNA tasks (see dnaBERT or dnaBERT2)
- DNA encodings share many similar traits to natural language
  - words, sentences, paragraphs
  - bases, codons, genes
- Generally
  - (1) Learn the “DNA language” through pretraining on reference genomes
  - (2) fine tune to particular tasks for inference



From Ji et. al. 2017




**NOAA**  
FISHERIES

# Transformer models are exceptionally good at paying attention to the whole sequence

## Example: Gene structure

- Because transformers use positional embeddings, it is trivial for the model to begin identifying “spatially important” signals:  
TATA boxes are ~30 bases upstream of ATG
- Increased learning depth could learn protein structure, taxonomic signal etc.




[https://en.wikipedia.org/wiki/TATA\\_box](https://en.wikipedia.org/wiki/TATA_box)

- Any function for which positional data is important could lend itself to the transformer model approach

# Usually, transformer models are trained in two steps


- **Pre-training**
- Large, unlabeled datasets
  - Wikipedia
  - RefSeq
- Self supervised
- Training task: Predict token based on surrounding context
- Learns the “language” (English, DNA, etc)
- Takes a **long** time
- **Fine-tuning**
- Smaller, labeled dataset
  - Coding sequences
  - Functional annotations
  - Taxonomy
- Model understands the language, so only needs to learn specific knowledge tailored to the task
- Rapid training
- **Inference**
- Using the model for predictive tasks



**NOAA**  
**FISHERIES**

# Overall: Transformers are uniquely able to natively handle DNA without transformation, which makes them powerful

- Cutting edge deep-learning approach
- Positional-awareness allows the model to learn spatial relationships
- Not limited by a kernel size
- Fine-tuning of pre-trained models allow for rapid application of old models to new tasks
- **Questions on Transformers?**



From:  
<https://www.exxactcorp.com/blog/Deep-Learning/a-deep-dive-into-the-transformer-architecture-the-development-of-transformer-models>